WHITE PAPER · ARCHITECTURE · PUBLIC DOMAIN

Intelligence at the Edge of the Cloud

A Local Architecture for the AI Era

Jonathan Horvat

2026 · NO RIGHTS RESERVED

ABSTRACT

The prevailing assumption in artificial intelligence is that intelligence lives in the cloud — that to be smart, a device must be connected. This paper argues the opposite will become true. As models mature and weight compression advances, the edge will not retreat into dumbness when disconnected. It will carry the full depth of human knowledge locally, refreshed not by a live wire but by periodic synchronization of model weights.

We propose a two-model edge architecture: a broadcast-synced sparse expert world model and an on-device augmented personal model with a built-in interpreter head. Combined, this architecture addresses three problems at once: intelligence that is genuinely personal, privacy that doesn't depend on trusting a corporation, and a fraction of the energy and infrastructure cost of cloud inference.

SECTION I

The False Promise of the Cloud

Every major AI deployment today rests on the same implicit bargain: give us your data, your queries, your context — and we will give you intelligence. The user sends their thoughts to a remote server. The server processes them against a vast model. The answer returns. The data stays.

This is not a design choice made for the user's benefit. It is a consequence of scale. The models are too large to run locally. The compute is too expensive to distribute. Every question you ask, every document you analyze, every private thought you process through an AI system passes through infrastructure owned by a handful of corporations.

> *The problem is not that cloud AI companies are untrustworthy. The problem is that trust should not be required.*

What follows is a description of an architecture in which it is not.

This is not an argument against the cloud. The cloud trains better models than any device ever will. It handles tasks that require computation at a scale no edge device can match. What this paper argues is simpler: not everything needs to be the cloud. A device that carries its own intelligence — that knows the world and knows you, locally, privately, offline — is not a retreat from capability. It is an addition to it. The architecture described here works alongside cloud AI, not instead of it. The local model is the default. The cloud remains available when you choose it. The difference is that choice now belongs to you.

SECTION II

Weights as Knowledge

A language model's weights are not code in the traditional sense. They are the accumulated result of exposure to human knowledge — billions of parameters shaped by contact with text, reasoning, and the patterns of thought that underlie both. A model's weights are, in a meaningful sense, a compression of what humanity knows — not a database, not a search index, but a form that encodes understanding rather than storing facts.

The sparse expert architecture described in this paper is not speculative. As of 2026, it is the dominant pattern at the frontier — the majority of leading model releases use sparse mixture-of-experts designs precisely because activating only a subset of experts per query delivers frontier capability at a fraction of the compute cost. The architecture inherits this efficiency directly.

Architectures improve, quantization and distillation advance, and the weight footprint for any given capability level continues to shrink. Capable models will run at the edge. The only question is when — and what the resulting architecture looks like.

→ [1]

SECTION III

The Sync Paradigm

A model weight update is categorically different from a software update. When a device receives updated weights, it does not merely gain new procedures. It absorbs new understanding. Months of collective learning compress into a payload the device integrates at rest, offline, without a persistent connection.

This is the sync paradigm: periodic, compressed, offline-compatible transfer of world-knowledge into edge devices. Not streaming. Not querying. Absorbing. In practice, a typical weekly sync payload for a curated expert subset is under 50 MB compressed — version-delta only. It fits comfortably in a background sync over Wi-Fi or metered 5G, with zero personal data ever leaving the device.

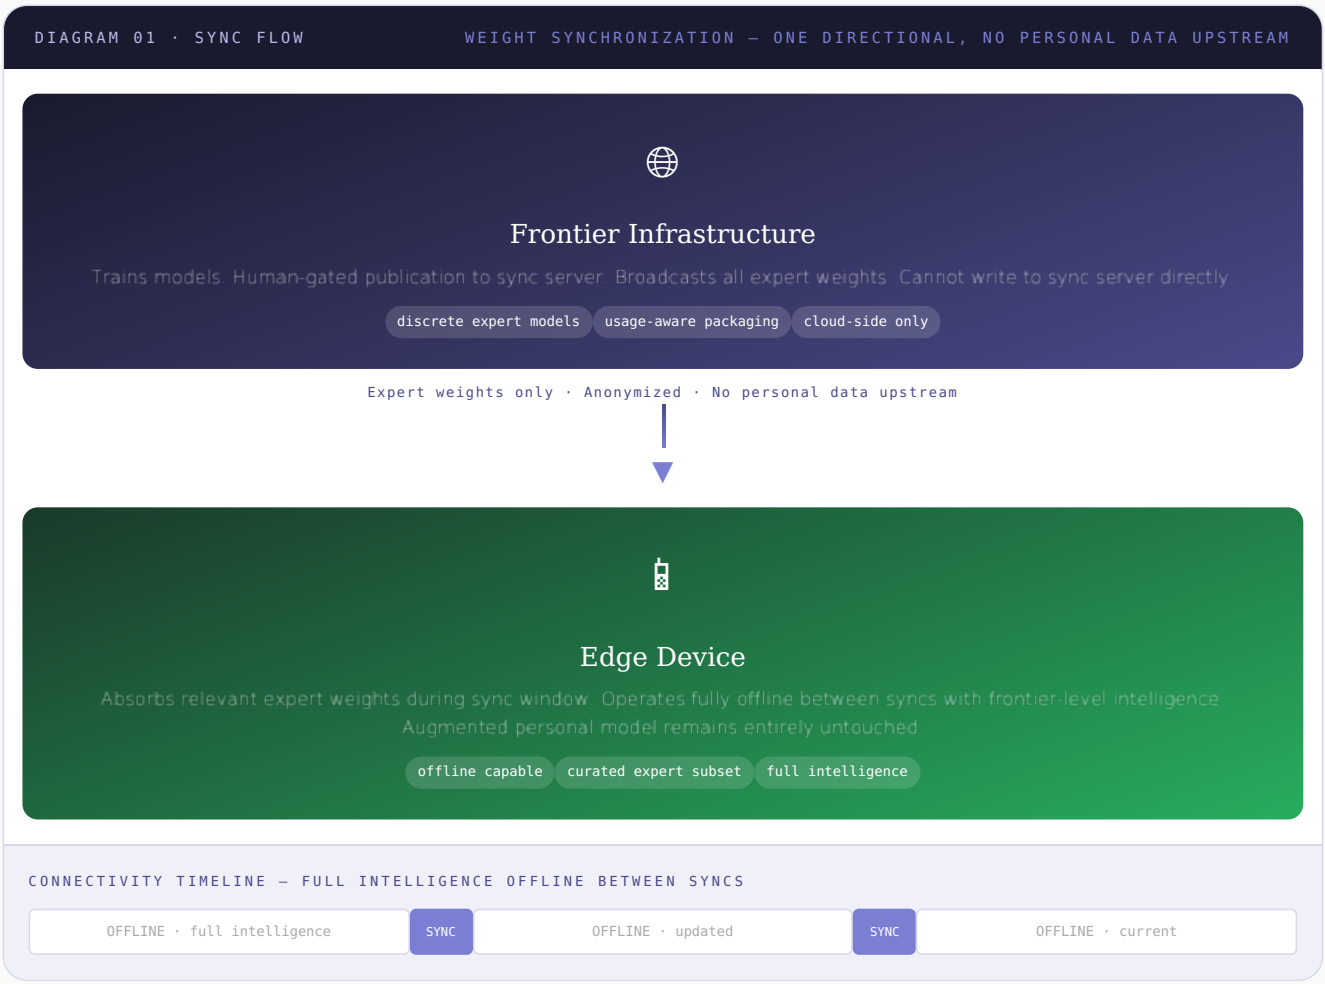

The sync request itself is the only fingerprint — and a remarkably thin one. It carries no query content, no personal context, no conversational history. With trivial hardening — batched anonymous requests through a proxy, timing randomization, CDN-style distribution of popular expert bundles — even this signal becomes indistinguishable from population-level noise. The privacy guarantee holds all the way down to the infrastructure layer.

> *Connectivity becomes optional. Intelligence does not.*

SECTION IV

The Core Architecture

Weight synchronization alone is not enough. A device that knows what humanity knows still lacks personal context.

DIAGRAM 02 · CORE ARCHITECTURE STACK

TWO CORE MODELS — WORLD MODEL AND AUGMENTED PERSONAL MODEL — ALL LOCAL

User

Query · intent · context — interfaces with augmented personal model only

receives from user ↓

AUGMENTED PERSONAL MODEL · SECURE ENCLAVE

Augmented Personal Model

Frozen base, trainable low-rank adapters encoding you, plus a fixed interpreter head for cross-space synthesis. Queries world model. Never transmits. Never exposes internal state.

frozen base · personal adapters · interpreter head · secure enclave · sole attack surface

query ↓ / output ↑

read/write ↕

WORLD MODEL · SPARSE EXPERTS · PULL-ONLY

World Model

Pull-only. Receives weight syncs. Responds to queries from augmented personal model. Cannot initiate, transmit, or access personal data.

pull-only · sparse expert routing · selective sync · no personal data

INTERPRETER HEAD — BUILT IN

Interpreter Head

Fixed, non-trainable cross-attention layer inside augmented personal model. Translates between independent embedding spaces. Frozen at compile time — never trained, never updated.

non-trainable · cross-attention · frozen at compile · inside secure enclave

COMPONENT	READS WORLD MODEL	ACCESSES PERSONAL ADAPTERS	REACHES CLOUD	RECEIVES FROM USER
User	—	—	—	—
Augmented Personal Model	Query only	Read / Write	—	Yes
World Model	—	—	Sync pull only	—
Interpreter Head	Output only	Read only	—	—

The World Model

A sparse expert model carries the distilled knowledge of the frontier. Individual expert models are maintained as discrete units — both at the cloud level and on device — enabling selective synchronization matched to usage profile. The world model is pull-only: it receives weight syncs and responds to queries from the augmented personal model. It cannot initiate communication, cannot transmit, and has no access to personal data.

The Augmented Personal Model

A small, efficient model runs continuously on-device atop a frozen base. It carries thin trainable low-rank adapter layers that encode you — your language patterns, your reasoning style, your domains of interest, your history.

Built directly into this model is a fixed, non-trainable interpreter head — a small cross-attention layer, frozen at compile time. This head is the only component that ever sees both the world model output and your personal adapter. It translates between the two independent embedding spaces without ever letting the models touch, then synthesizes the final response.

The entire personal delta — adapters plus interpreter head — remains just a few megabytes. It is updated incrementally through interaction and never leaves the device. The interpreter head lives inside the same hardware secure enclave as the personal model, making the attack surface extremely narrow by design.

The user interacts only with the augmented personal model. The world model never sees personal context. The personal model never exposes its internal state. All orchestration happens locally, inside a single secure component.

```
→ [2, 3, 4]
```

SECTION V

Privacy as Architecture, Not Policy

DIAGRAM 03 · PRIVACY THREAT MODEL

ATTACK SURFACE ANALYSIS — CLOUD MODEL VS EDGE ARCHITECTURE

⚠ CURRENT CLOUD MODEL

Personal data leaves device
Every query, document, and context transmitted to remote servers just to receive a response.

Corporate trust required
Privacy is a policy promise — revisable, breachable, erodable by business model changes.

Large diffuse attack surface
Data in transit, at rest on servers, in logs — multiple external exposure points.

✓ EDGE ARCHITECTURE

Augmented personal model never leaves device
Local by design. No query requires transmitting personal context.

Privacy is structural
A consequence of the architecture. Cannot be revoked by policy or eroded by business model.

Single narrow attack surface
Only the augmented personal model touches both — via its built-in interpreter head — locally, inside a secure enclave.

Broadcast eliminates the fingerprint
No request-response means no fingerprint. The server broadcasts to everyone. It never knows who kept what.

⚠ RESIDUAL ATTACK SURFACE

The **augmented personal model** is the sole component requiring hardening. **Secure enclaves** and **trusted execution environments** — already mature in consumer hardware — map directly onto this requirement. Best practices such as side-channel resistance, regular firmware updates, and runtime attestation further harden this single narrow surface. One local target is categorically more defensible than distributed cloud infrastructure.

DIAGRAM 04 · QUERY FLOW COMPARISON

CLOUD PARADIGM VS EDGE ARCHITECTURE

Cloud Paradigm

ORIGIN
User forms a query

↓

⚠ EXPOSURE
Personal context transmitted to remote server

↓

⚠ PROCESSING
Processed against cloud model — data at rest on external infrastructure

↓

⚠ DEPENDENCY
No connectivity = no intelligence

↓

RESULT
Intelligent response returned

VS

Edge Architecture

ORIGIN
User forms a query

↓

✓ LOCAL
Augmented personal model receives intent — nothing leaves device

↓

✓ LOCAL
World model queried · Interpreter head synthesizes · All local

↓

✓ RESILIENCE
Works fully offline — intelligence is a local property

↓

RESULT
Intelligent response — private, local, current

> *Privacy becomes a property of the system, not a feature of the product. It cannot be revoked by a policy update. It cannot be eroded by a business model change. It simply is.*

SECTION VI

Graceful Degradation as a First-Class Feature

A device with limited or outdated experts is not broken — it simply has known, expressible limits. The augmented personal model tracks exactly which weights are present, how current they are, and how well-matched they are to the query at hand.

Rather than hallucinating into gaps, the interpreter head surfaces a coverage score — a weighted combination of router activation confidence, expert freshness relative to the query domain, and embedding-space domain relevance. The result is a concrete honest signal: "*Drawing from world knowledge current to March 2026 with 92% domain coverage for quantum error correction.*"

A user asks about the latest developments in quantum error correction given their prior work on surface codes. The system returns the coverage score, leans on personal adapters to tailor the response to their exact reasoning style and historical context, and surfaces uncertainty where it exists. No hallucination — just transparent, graceful degradation.

In degraded states it leans more heavily on what it knows about the user to fill gaps intelligently. The personal adapters don't degrade with the world model. They always reflect you.

SECTION VII

Neural Architecture

Two core models, each with independent embedding spaces, coordinated by a built-in interpreter head inside the augmented personal model.

Interpreter Head Mechanics

The interpreter head is a small, fixed, non-trainable cross-attention module — typically 4–8 heads, low dimensionality — frozen at compile time. It takes the world model's output tokens in their native embedding space and attends to them from the personal model's context tokens. Because it is frozen and read-only, no gradients ever flow back into either model during adapter training. It acts purely as a translator while preserving strict separation.

```
world_output  = world_model(query_from_personal)    # embedding space W
personal_ctx  = personal_adapters(user_history)     # embedding space P
translated    = interpreter_head.cross_attention(
    query=personal_ctx,
    key=world_output,
    value=world_output
)
final_response = personal_model.decode(translated)
```

The world model remains a sparse expert system that is pull-only and periodically refreshed via sync. The augmented personal model contains both the trainable low-rank adapters that encode the user and a fixed interpreter head that performs the cross-space synthesis. This head bears the full computational cost of translation while preserving strict separation: the world model never sees personal data, and the personal model never leaks its private state outward.

All other architectural benefits — selective sync, graceful degradation, per-expert quantization, and incremental on-device training — remain unchanged.

Neither component requires new invention. Sparse expert world models are already the dominant frontier architecture. Adapter-based on-device personalization with a frozen base is already in production — major device manufacturers ship exactly this pattern today. The contribution here is the combination: privacy-first, offline-resilient, locally intelligent.

DIAGRAM 05 · NEURAL ARCHITECTURE BLUEPRINT

TWO CORE MODELS — INTERPRETER HEAD BUILT INTO AUGMENTED PERSONAL MODEL

WORLD MODEL

SPARSE EXPERTS · BROADCAST RECEIVE

BASE
Shared Router
Routes queries to relevant experts. Frozen between syncs.

EXPERT A
Domain Expert
Higher precision. Frequently used. Prioritized in sync.

EXPERT B–N
Domain Experts
Lower precision fallback. Synced by usage profile.

QUANTIZATION
Per-expert precision
Frequent experts at higher precision. Fallbacks aggressively quantized. Profile evolves with usage.

EMBEDDING
Independent space
Does not share embedding space with augmented personal model.

Broadcast: Server broadcasts all expert weights. Devices retain what they use locally. Bandwidth cost is version-delta only — typically <50 MB/week compressed.

↕ queried by augmented personal model only · through user intent as lens

AUGMENTED PERSONAL MODEL

SECURE ENCLAVE · INTERPRETER HEAD BUILT IN

INTERPRETER HEAD
Fixed cross-attention
Non-trainable. Frozen at compile time. Translates between world model output and personal adapter space. Bears full cross-space synthesis cost.

PERSONAL ADAPTERS
Low-rank delta
Trainable layers encoding you. A few MB. Updated by behavioral observation only. Never transmitted. Never supervised.

FROZEN BASE
Shared foundation
Shared across interpreter head and personal adapters. Broadcast on slow cadence. Geometry-constrained updates preserve adapter integrity.

DEGRADATION
Coverage score
Tracks expert presence and freshness. Surfaces coverage score rather than hallucinating. Personal adapters always reflect you regardless of world model state.

Security: Runs inside a hardware secure enclave. The sole attack surface of the entire architecture. Dedicated hardware-level protection is a fundamental architectural requirement.

→ [4, 7, 8]

An Open Research Direction

The interpreter relationship described here may itself be more than a fixed mechanism — it may be emergent. Prior work in multi-agent reinforcement learning has shown that neural networks develop compressed, task-specific communication protocols unprompted when given a consistent interface and a signal that rewards successful coordination. The personal model learning to read the world model through observation may follow the same principle: given enough interaction, a high-bandwidth AI-native channel develops that is more efficient than anything a human designer would specify, and specific to this device, this user, this pair of models.

Two findings from the broader literature suggest this channel may be more learnable than expected. Large language models have been shown to develop functional organization strikingly similar to human brain fMRI activation patterns — consistent specialization emerging independently across models trained on similar architectures. And the human brain itself exhibits fractal organizational structure across scales, with growing evidence that sufficiently trained artificial networks converge toward similar self-similar structure under the pressure of efficient compression. If independently trained models develop analogous internal organization, the personal model is not learning to read an alien space — it is finding correspondence between two instances of the same underlying pattern.

Whether oscillatory synchronization between models — analogous to neural entrainment in biological systems — could accelerate the formation of this channel, and whether fractal scale-invariance makes it transferable across capability upgrades, are open research questions. The architecture creates the conditions. What emerges is worth measuring.

→ [5, 6, 9, 10, 11, 12, 13, 14, 15]

SECTION VIII

The Graduated Transition

This architecture does not require a utopian leap. It is viable today at reduced capability, and becomes more capable as hardware advances. The weight complexity synced to a device scales with what that device can run. The architecture is constant across the capability curve — only the ceiling changes.

As edge chips improve, larger and more capable world model weights become syncable. As world models become more efficient through distillation and quantization, they fit on less capable hardware. Both curves accelerate each other. The transition is a ramp that builds itself.

In a sparse expert architecture, different expert weights can carry different precision profiles — frequently used experts at higher precision, rarely accessed fallbacks more aggressively quantized. The augmented personal model's usage patterns inform this profile over time — the device becomes progressively better calibrated to its owner without any external direction.

Today, a mid-range smartphone running a 3–7 billion parameter quantized world model delivers genuine general intelligence locally — summarization, reasoning, writing, analysis — across the domains its expert profile covers. The personal adapter set starts empty and becomes calibrated within days of normal use. The architecture is useful from day one and better by day thirty. At the other end of the curve, as dedicated edge AI chips mature and model compression continues, the gap between local and cloud capability narrows to the point where the distinction becomes about latency and privacy preference rather than raw intelligence. The transition does not require waiting for that end state. Throughout it, the cloud remains available for tasks that exceed local capability. Early adopters route to the cloud selectively. Later adopters rarely need to. The architecture accommodates both without changing.

SECTION IX

Physical AI

The cloud model is not merely inefficient for physical AI. It is architecturally incompatible with it.

01 Local intelligence is a hard constraint

Physical systems operate in real time. When the consequence of latency is physical — a fall, a collision, a missed action — a cloud round trip is not slower, it is disqualifying. Intelligence must live where the action is.

02 Connectivity in physical environments is unreliable by nature

Physical environments are dynamic, adversarial, and unpredictable. An architecture premised on connectivity will fail at the moments that matter most. An architecture premised on local intelligence with periodic sync will not.

03 Physical systems that know their operator are more capable and safer

A system calibrated to a specific person — their patterns, their tolerances, their preferences — performs better and fails more gracefully than a generic one. The personal adapter model achieves this privately, without transmitting anything about the person.

04 The broadcast model scales naturally to physical deployment

Fleets, networks, distributed systems — all receive the same broadcast, all operate independently, all stay current without central dependency or single point of failure. There is no architecture to coordinate, no hub to lose.

05 Conservative failure is a safety property

A cloud-dependent physical system that loses connectivity faces two bad options: stop, or continue on intelligence of unknown currency. A local system with explicit uncertainty knows exactly how current its world model is and can fail conservatively. Graceful degradation in physical AI is not a convenience feature — it is a design requirement.

This architecture does not just accommodate physical AI. It is the architecture physical AI requires.

→ [18, 19, 20]

SECTION X

Consequences at Scale

Infrastructure Displacement

The current AI infrastructure buildout is premised on inference at scale — serving billions of queries continuously through cloud architecture. The ongoing energy cost of AI lives here, not in training. Moving inference to the edge eliminates this cost at the cloud level. The aggregate power consumption of edge devices is not zero — but the efficiency advantage of local inference over remote inference is so large that the net reduction is substantial. The ongoing energy cost of AI becomes a bounded problem rather than an open-ended one.

Beyond energy, the continuous internet infrastructure that exists primarily to shuttle queries to models and return answers becomes largely unnecessary for AI interaction. The infrastructure problem becomes one of occasional high-quality sync windows rather than permanent fat pipes — a categorically different and cheaper problem.

Geopolitical Leapfrogging

The mobile phone leapfrogged landline infrastructure across much of the developing world. This architecture enables a more complete leapfrog for AI and, ultimately, general connectivity. The only infrastructure requirement is the device itself and an occasional low-bandwidth sync window — deliverable by satellite, intermittent wifi, or peer-to-peer mesh networks sharing weight updates between devices.

Nations that have resisted cloud AI adoption because it meant routing citizens' data through foreign infrastructure gain an alternative that requires no such compromise. Intelligence becomes a local resource rather than a foreign service. A government can cut internet access. It cannot cut intelligence that is already local.

The Edge Chip as Strategic Battleground

If inference moves to the edge permanently, the semiconductor industry's priorities shift. The device chip must run world model inference and the augmented personal model continuously, efficiently, and securely. Power efficiency, unified memory, neural processing unit integration, and hardware-level secure enclaves become the metrics that matter.

A new measure emerges to replace tokens-per-second: intelligence per milliwatt. Whoever defines and wins that metric defines the next era of semiconductors. Hardware secure enclaves for the augmented personal model become a durable competitive moat.

→ [16, 17]

SECTION XI

The Human Analogy

The architecture maps naturally onto human cognition. The personal adapters correspond to episodic and autobiographical memory — individual, never shared wholesale. The world model corresponds to semantic memory — general knowledge of how things work. The augmented personal model corresponds to executive function — dynamically combining both in context.

These systems in the human mind do not merge. They are accessed and combined dynamically, in response to context. The architecture didn't invent this separation — it encodes something cognition already figured out.

SECTION XII

Environmental Implications

The prevailing assumption in AI infrastructure is that intelligence scales with physical build-out. More compute means more capability. More capability means more data centers. More data centers mean more cooling, more power, more fiber, more buildings, more land, more water. This is not a temporary phase. It is the foundational logic of the cloud model — and its environmental consequences are structural outputs of that model, not incidental side effects.

This architecture inverts that assumption. Intelligence stops being something you build out and starts being something you compress inward.

> *The cloud model treats intelligence as a utility that must be generated centrally and transmitted. This architecture treats intelligence as a property that can be stored locally and updated periodically. The environmental difference between those two models is the difference between a power grid and a battery.*

Energy — Inference, Not Training

The public conversation about AI's energy footprint focuses heavily on training — the large, visible, one-time cost of building a model. But training is bounded and periodic. Inference is continuous, and it scales with every user on earth, every query, every day. Moving inference to the edge eliminates this ongoing cost at the cloud level. Training infrastructure remains — but training a model once and distributing its weights is categorically different from serving that model's outputs billions of times daily.

Water

Data centers are large consumers of water, used primarily for cooling — a cost largely invisible in public accounting of AI's environmental impact. The water footprint of a single large data center can rival that of a small city. As construction accelerates globally, this burden falls on local watersheds and aquifers, often in regions already under stress, chosen precisely because land and power are cheap.

Geographic Concentration of Burden

Cloud infrastructure does not distribute its environmental burden evenly. Data centers cluster near cheap power and water — which means specific regions, specific communities, and specific ecosystems absorb disproportionate costs. Rivers near data center clusters run warmer. Local grids are stressed. Distributed edge intelligence diffuses this burden across billions of devices that people are already powering. No community bears a concentrated environmental cost. No watershed is selected for proximity to cheap land.

Intelligence Per Milliwatt

The current dominant metric in AI hardware is tokens per second — optimized for cloud-scale inference serving. For the edge architecture, the right metric is different: intelligence per milliwatt — how much genuine reasoning capability can be delivered per unit of energy on a constrained device. This reorientation builds energy efficiency into the fundamental design objective of the next generation of AI hardware, rather than treating it as a secondary concern.

> *The history of energy suggests that distributed storage is more resilient, more equitable, and ultimately more efficient than centralized generation and transmission at scale. The same logic applies to intelligence.*

```
→ [16, 17]
```

SECTION XIII

On Sharing This Freely

This paper is public domain. The ideas here belong to no one. They are offered to the commons in the same spirit that the foundational architectures of the internet were offered — not as products to be extracted from, but as gifts to a future that will build on them in ways their originators could not predict.

If something here is useful, use it. If something here is wrong, correct it publicly. If something here sparks a better idea, share that too.

SECTION XIV

Reference Implementation Sketch

This architecture is implementable today on any device with a neural processing unit and hardware secure enclave — Apple A/M-series with Secure Enclave, Qualcomm Snapdragon with Hexagon NPU and TrustZone, and equivalents. The world model runs on existing sparse-MoE runtimes. Personal adapters use on-device LoRA training loops. The interpreter head ships as a static frozen ONNX or CoreML module. Sync logic is a background asset pipeline with delta patching.

No new silicon or OS primitives are required. The architecture composes what already ships in 2026 consumer hardware. The only thing missing is the decision to build it this way.

REFERENCES

[1] Dai, D. et al. (2024). *DeepSeekMoE: Towards Ultimate Expert Specialization in Mixture-of-Experts Language Models.* arXiv:2401.06066.

[2] Hu, E. et al. (2021). *LoRA: Low-Rank Adaptation of Large Language Models.* arXiv:2106.09685.

[3] Kopiczko, D. J. et al. (2025). *Low-rank adaptation for edge AI.* Scientific Reports.

[4] Apple Machine Learning Research (2024/2025). *Apple Intelligence Foundation Language Models.* machinelearning.apple.com. — Adapter-based on-device personalization with frozen base, in production across Apple devices.

[5] Liu, Y. et al. (2025). *Brain-Inspired Exploration of Functional Networks and Key Neurons in Large Language Models.* arXiv:2502.20408.

[6] AlKhamissi, B. et al. (2024). *The LLM Language Network: A Neuroscientific Approach for Identifying Causally Task-Relevant Units.* arXiv:2411.02280.

[7] Anthropic (2023). *Towards Monosemanticity: Decomposing Language Models With Dictionary Learning.* transformer-circuits.pub.

[8] Bereska, L. & Gavves, S. (2024). *Mechanistic Interpretability for AI Safety — A Review.* arXiv:2404.14082.

[9] Wang, Y. et al. (2024). *Neuro-evolutionary evidence for a universal fractal primate brain shape.* eLife. DOI:10.7554/eLife.92080.

[10] Hilgetag, C. C. & Goulas, A. (2020). *The fractal brain: scale-invariance in structure and dynamics.* Cerebral Cortex, 33(8), 4574–4587.

[11] Miyato, T., Löwe, S., Geiger, A. & Welling, M. (2025). *Artificial Kuramoto Oscillatory Neurons (AKOrN).* ICLR 2025 (Oral). arXiv:2410.13821.

[12] Shim, G. et al. (2025). *HoloBrain and HoloGraph: Brain-inspired oscillatory synchronization for graph neural networks.* Nature Communications.

[13] Scientific Reports (2025). *Deep Oscillatory Neural Network (DONN).* Vol. 15, Article 40968.

[14] Zhu, C. et al. (2024). *A Survey of Multi-Agent Deep Reinforcement Learning with Communication.* Autonomous Agents and Multi-Agent Systems, Vol. 38.

[15] Li, Y. et al. (2015). *Convergent Learning: Do Different Neural Networks Learn the Same Representations?* arXiv:1511.07543. — Empirical evidence that independently trained networks converge on similar representational features.

[16] MIT Technology Review (2025). *We did the math on AI's energy footprint.* technologyreview.com. — Inference now accounts for 80–90% of AI computing energy.

[17] International Energy Agency (2024). *Energy and AI: Energy Demand from AI.* iea.org. — Global data center electricity consumption projected to double to 945 TWh by 2030, with AI as the primary driver.

[18] Chafii, M. et al. (2023). *Emergent Communication in Multi-Agent Reinforcement Learning for Future Wireless Networks.* arXiv:2309.06021.

[19] TechRxiv (2025). *Physical AI: Bridging the Sim-to-Real Divide Toward Embodied Intelligence.* — Latency, bandwidth, and reliability constraints render cloud-centric inference architectures insufficient for physical AI.

[20] Edge AI and Vision Alliance (2026). *On-Device LLMs in 2026: What Changed, What Matters, What's Next.* — Cloud round-trips add hundreds of milliseconds, breaking real-time physical experiences.

Public domain. No rights reserved. Build freely.

JONATHAN HORVAT